

VC-Funded RL Environment / Task / Verifier Companies (2022-2026)

VC-Funded RL Environment / Task / Verifier Companies (2022 – Jul 2026)

Exhaustive market inventory. Two-pass method (run 2026-07-02): **8 independent discovery lenses** (news, VC portfolios, funding databases, analyst maps, accelerators, global, verifiers-specific, rebrands/M&A) → **59 new companies** on top of a ~20-company seed → **76 enriched** for funding-in-window, scope, tier, status, and a one-sentence differentiation claim.

Scope: companies building (a) RL **environments**/gyms, (b) agentic **tasks**/data-for-RL, or (c) **verifiers**/reward-models/RLVR — for LLM/agent training. **Excluded per your instruction:** robotics/AV/embodied-RL/physical simulation. Global. Accelerator-stage included; bootstrapped listed separately.

Result: 60 qualifying companies (in-scope, VC/accelerator-funded, round in-window) across 3 tiers; 2 bootstrapped/public; 11 excluded (out of scope). ^{a c} = since acquired. Full funding rounds, founders, sources per company in `rl-env-funded-companies.csv`.

Caveats: many are seed/pre-seed with **undisclosed** amounts (YC/accelerator deals); several figures come from aggregators (Crunchbase/Tracxn/StartupHub) not primary filings — confidence is per-row in the CSV. Same-name disambiguation was applied throughout (Refresh, Lucid, Theta, Dojo, Matrices, Composio are common names); notable identity flags are in each row's redFlags.

Tier 1 — Pure-plays (core product is RL environments / tasks / verifiers) — 35

Company	Cat	Geo	Total raised	Key round	Differentiation (their claim)
Andon Labs	mixed	San Francisco	~\$500K	YC Winter 2024 batch (pivoted...	Builds long-horizon agent benchmarks/environments (Vending-Bench, Butter-Bench, Blueprint-Bench) that measure agentic coherence and safety, and runs ...
Andromede	mixed	Lausanne	Undisclosed	Backed by Unusual Ventures (s...	RL data lab that programmatically generates environments, tasks, and verifiers from real-world data for post-training and evaluation of frontier agen...
Aviro	mixed	San Francisco	Undisclosed (YC...	Pre-seed round (amount, lead ...	Builds RL environments and benchmarks (e.g. Enterprise Search Benchmark) for long-horizon tool use across ML research, live web, and enterprise knowl...
Axiom Math	verif	US (Bay Area ...	\$64M	Seed: 64Mat 300M valuatio...	Builds a self-play 'conjecturer + prover' loop that generates and mechanically proves math problems, outputting machine-checkable Lean proofs (AxiomP...
BenchFlow	mixed	San Francisco	~\$1M (seed, rep...	Seed round ~\$1M (per aggregat...	Open-source 'environment lab' providing a universal framework to run agents against RL environments (SkillsBench, ClawsBench mock-workplace envs) plu...
Bespoke Labs	mixed	Mountain View	\$40M	\$40M total announced July 6, ...	Pairs world-class data-curation research (OpenThoughts dataset, Bespoke Curator, Terminal-Bench) with best-in-class RL environments to train and veri...
Chakra Labs	env	Brooklyn	~\$10.1M	Seed round ~\$10.1M, reported ...	Builds high-fidelity, deterministic computer-use RL environments (SPA/app clones) and human trajectory datasets — its Dojo suite brings frontier-lab-...
Collinear AI	mixed	San Francisco	Undisclosed (se...	Seed round (undisclosed amoun...	SimLab builds sandboxed, stateful enterprise RL environments (simulating tools like Jira, ServiceNow, Shopify, EMR, airline/hotel systems) that emit ...
Datacurve	mixed	San Francisco	~\$17.7M	Series A: \$15M, announced Oct...	An expert-curated coding data foundry (gamified 'Shipd' bounty platform with 14,000+ vetted programmers, paying for outputs not hours) supplying fron...
Deeptune ^{ac}	env	New York	\$43M (pre-acqui...	Series A: \$43M, closed/announ...	Builds high-fidelity RL 'training gyms' that simulate real professional workflows (accounting, customer support) across apps like Slack and Salesforc...
Dojo (by Chakra Labs)	env	Brooklyn	N/A (rolls up t...	No standalone financing — Doj...	A collaborative/open RL-environment suite giving computer-use agents deterministic, frame-accurate clones of production software plus crowd-sourced h...
Fleet	env	San Francisco	~\$15M closed se...	Series A REPORTED/CLOSING Apr...	Builds RL training 'gyms' — high-fidelity simulated replicas of enterprise software like Salesforce and Excel — that frontier labs pay to train agent...
General Reasoning	mixed	London	~\$10.9M	Seed: \$10,904,992 in equity p...	Builds open reasoning-data and RL-environment infrastructure — reasoning datasets and verifiable-reward (RLVR) tooling/environments — for frontier fo...
Habitat	env	New York	Undisclosed	A Crunchbase seed funding-rou...	Builds RL environments for white-collar / agentic work automation — code and desktop (computer-use) interaction environments offering hundreds of div...
Haize Labs	verif	New York City	~\$12.5M	Seed: \$12.5M led by General C...	Builds Verdict, a declarative framework for composing scalable LLM judges / automated verifiers and reward models, paired with large-scale automated ...

Company	Cat	Geo	Total raised	Key round	Differentiation (their claim)
Halluminate	mixed	San Francisco	Undisclosed (ag...	Y Combinator S25 (Summer 2025...	Builds computer/browser-use RL gyms — a fully-simulated internet ('Westworld') of synthetic consumer/enterprise apps — with verifiable rewards, narro...
Harmonic AI	verif	Palo Alto	~\$295M+	Series C: \$120M at \$1.45B val...	Builds 'Mathematical Superintelligence' via its Aristotle model that translates natural-language math into Lean 4 and formally verifies every proof s...
HUD	mixed	San Francisco	Undisclosed (YC...	Seed (undisclosed amount, ~20...	Open-source SDK + hosted gateway to wrap real software as isolated, agent-callable RL environments and evals specifically for computer-use and browse...
Idler	env	San Francisco	Undisclosed (YC...	Pre-seed / accelerator: Y Com...	Builds reinforcement-learning training environments from real-world, expert-level coding problems so foundation labs can train and evaluate code-gene...
Judgment Labs	verif	San Francisco	\$32M	\$32M total across a Seed and ...	Builds a 'continuous improvement layer' / evaluation-and-verifier infrastructure for deep AI agents that scores long reasoning traces, tool use and m...
Lucidic AI	mixed	San Francisco	~\$500K	YC Winter 2025 (W25) batch; \$...	A training/eval platform that ingests real agent logs and uses controlled simulations plus RL, genetic algorithms and Bayesian optimization — scored ...
Maingen	env	San Francisco	Undisclosed (YC...	Y Combinator S26 accelerator ...	Builds RL environments for industrial operations (framed as a ~\$5T slice of the US economy) so frontier labs can train models that run factories.
Mechanize	mixed	San Francisco	~\$9.1M	Seed: \$9.1M at \$500M post-mon...	Builds robust, high-fidelity RL environments, benchmarks, and training data ('digital office' simulations) for frontier coding/agentic tasks, working...
Micro1	task	San Francisco...	~\$41.6M	Series A \$35M led by 01 Advis...	Runs an AI-vetted, top-1% expert network (PhDs, senior engineers) supplying on-demand RLHF and agentic training data to frontier AI labs, positioned ...
Plato	env	San Francisco	Undisclosed (pr...	Pre-seed/seed (2025) — invest...	Builds simulated 'worlds' — high-fidelity replicas of real websites/software (Amazon/Airbnb/Gmail-style) plus a full Linux desktop 'Computer Use' tar...
Preference Model	env	San Francisco	Undisclosed	Seed round (amount undisclose...	Builds high-quality RL training environments that reflect real-world complexity with diverse tasks and robust reward functions, partnering directly w...
Proximal	mixed	San Francisco	Undisclosed (se...	Seed round (amount undisclose...	Builds high-fidelity, long-horizon RL coding environments grounded in real codebases using software-driven data engines (not human contractors), plus...
Refresh	mixed	San Francisco	Undisclosed (YC...	Y Combinator X25 (later relab...	Builds simulation-environment 'training gyms' with deterministic and rubric-based verifiable rewards that turn real software-engineering and computer...
Rubric AI	verif	San Francisco	~\$500K (reporte...	Seed / YC standard deal — par...	Turns credentialed expert judgment into rubric-based reward signals and RL environments for non-verifiable/high-stakes domains where ground-truth ver...
Sepal AI ^{ac}	mixed	San Francisco	~\$500K (pre-acq...	YC Summer 2024 batch; ~\$500K ...	Data-research company producing high-quality training data, expert-graded evaluation benchmarks and RL environments for frontier LLMs, drawing on a n...
Taste Labs	mixed	New York	\$18.5M	Seed: \$18.5M, June 2026, co-l...	Builds the 'taste layer' — preference datasets, rubrics, evaluation environments and RL post-training data to give frontier models and agents subject...

Company	Cat	Geo	Total raised	Key round	Differentiation (their claim)
The LLM Data Company (TLDC / doteval)	mixed	San Francisco	500K–3.6M (co...	YC X25 (Spring 2025) + pre-se...	Builds post-training data and RL environments for frontier labs in critical/expert domains, plus 'doteval' tooling to write, version and execute eval...
Theta (Theta Software)	env	San Francisco	Undisclosed (YC...	Y Combinator X25 (Spring 2025...	Translates real-world enterprise workspaces/software into high-fidelity simulation environments to train and post-train expert-level computer/browser...
Veris AI	env	San Francisco	\$8.5M	Seed: \$8.5M, closed/announced...	Trains enterprise AI agents through high-fidelity simulated experience (RL-style learning by doing) rather than prompt engineering or human-labeled d...
Vmax	mixed	San Francisco	Undisclosed (ea...	Backed by Race Capital and So...	Automates converting a customer's proprietary data and evals into RL environments for LLM-based agents, targeting long-horizon and coding tasks (e.g....

Tier 2 — Incumbent data/eval companies with an RL-environments line — 7

Company	Cat	Geo	Total raised	Key round	Differentiation (their claim)
Alignerr (Labelbox)	task	San Francisco	~\$189M (parent ...	Alignerr is not separately fi...	Positions itself as an on-demand marketplace of 2.6M+ vetted domain experts generating RLHF, complex-reasoning and multimodal reward signals for fron...
Centific	env	Redmond / Sea...	\$60M	Series A: \$60M, June 2025, le...	Launched 'RL Environments-as-a-Service' — configurable, industry-authentic simulated enterprises for training and evaluating enterprise AI agents aga...
Huzzle Labs	mixed	London	~\$6M (reported)	~\$6M raised (aggregator-repor...	Human-intelligence 'data foundry' bundling long-horizon RL environments, expert trajectory data, and contextual evals for code/tool-use/computer-use/...
Invisible Technologies	env	New York	~\$144M	\$100M growth round, Sep 16 20...	Builds production-grade RL environments from real enterprise workflows with domain-expert-designed reward functions and fully logged, inspectable, re...
Pareto AI	task	San Francisco	~\$14M (aggregat...	Seed: \$4.5M closed 14 Mar 202...	Talent-first expert human-data platform positioning itself as 'the verification layer for reinforcement learning on real-world expertise,' supplying ...
SuperAnnotate	env	San Francisco	~\$50M+	Series B: \$36M, Nov 2024, led...	Data-annotation/data-management incumbent now designing and running end-to-end RL environment programs — collecting human trajectories and building r...
Toloka	mixed	Amsterdam	\$72M disclosed ...	May 7, 2025: \$72M (~€64M) str...	A large established AI-data/eval provider that has added RL-gyms and multi-user virtual-organization environments producing high-fidelity trajectorye...

Tier 3 — Adjacent infrastructure (RL training infra / agent sandboxes; in the ecosystem, not a core env/task/verifier product) — 18

Included for completeness because they're frequently cited in this market, but they sell training infra or sandboxes rather than environments/tasks/verifiers as the product.

Company	Cat	Geo	Total raised	Key round	Differentiation (their claim)
Adaptive ML	verif	New York	\$20M	\$20M seed (announced March 12...	Adaptive Engine is an enterprise RLOps platform that fine-tunes, evaluates and serves open-source LLMs using RLHF/RLAIF, generating reward signals an...
AgileRL	mixed	London	~\$7.5M	Seed: £5.5M / ~\$7.5M, announc...	Open-source RL framework plus a managed full-stack RLOps product (Arena) that speeds up reinforcement learning ~10x with distributed training and env...
Composio	env	San Francisco	~\$29M	Series A: \$25M led by Lightsp...	Positions as the 'learning layer' for agentic AI — providing tool-calling infrastructure (200+ app integrations) plus RL-style environments and skill...
Coval	env	San Francisco	~\$31M	Series A \$28M led by Norwest ...	Applies autonomous-vehicle-style simulation (millions of simulated test runs) to stress-test, evaluate and monitor voice and chat AI agents for relia...
E2B	env	San Francisco	~\$32M (aggregat...	Series A: \$21M led by Insight...	Provides secure, isolated cloud code-execution sandboxes as the runtime substrate for AI agents — increasingly adopted as the execution environment l...
Gray Swan AI	verif	Pittsburgh	~\$45M+ (Series ...	Series A: \$40M, closed ~May/J...	AI-security company whose Arena red-teaming community (15,000+ researchers) and Shade adversarial-testing engine generate the adversarial data and gr...
Kaizen (Kaizen Automation)	env	San Francisco	~\$4M+	Seed: 'over \$4M' (reported ~O...	Builds RL environments that simulate real operational 'work' — legacy web-portal/back-office tasks across healthcare, logistics, and financial servic...
Maxim AI	verif	San Francisco	\$3M	Seed: \$3M, announced Jun 2024...	End-to-end agent simulation, evaluation and observability platform with a configurable graders/evaluators engine to test AI agents across scenarios w...
Morph Labs	mixed	San Francisco	~\$5.75M	Seed: \$5.75M led by Khosla Ve...	Builds Infinibranch, a rapid compute-branching / snapshot environment infrastructure for massively parallel agents, aimed at autoformalization and 'v...
Nous Research	env	New York / Un...	~65M(≈15M ea...	Series A \$50M led by Paradigm...	Publishes Atropos, an open-source LLM reinforcement-learning environments framework (1,200+ tasks across math/code/tool-use) for collecting and evalu...
OpenPipe ^{ac}	verif	Seattle / Bel...	\$6.7M disclosed...	Acquired by CoreWeave (announ...	Open-source Agent Reinforcement Trainer (ART) that RL-fine-tunes agents via rollouts, custom reward functions (RULER), and GRPO to teach agents to re...
Osmosis	mixed	San Francisco	\$6.3M	Seed: \$6.3M, announced Oct 20...	A forward-deployed reinforcement fine-tuning platform that captures customers' real-world feedback, converts it into structured reward functions, and...
Periodic Labs	verif	San Francisco	\$300M (confirme...	Seed: \$300M, announced Sept 3...	Building an autonomous physical lab so that real experiments serve as the RL environment and produce experimentally-verified reward signals ('nature ...
Predibase ^{ac}	verif	San Francisco	~\$28M pre-acqui...	Acquired by Rubrik, announced...	End-to-end reinforcement fine-tuning (RFT) platform that lets teams fine-tune open-source models using user-defined reward functions/graders instead ...
Prime Intellect	mixed	San Francisco	>\$150M	Seed extension: \$15M in Feb 2...	Runs an open 'Environments Hub' ('Hugging Face for RL environments', 2,500+ community environments plus verifiers and the PRIME-RL framework) alongsi...
Runloop	env	San Francisco	\$7M	Seed: \$7M, closed/announced J...	Provides enterprise-grade 'Devboxes' — secure isolated micro-VM sandboxes plus GitHub integration, snapshots, blueprints, and public coding benchmark...
RunRL	mixed	San Francisco	~\$500K	Seed / accelerator round ~\$50...	RL-as-a-service platform where users define a metric/reward and RunRL trains their model or agent with reinforcement learning without them managing G...

Company	Cat	Geo	Total raised	Key round	Differentiation (their claim)
TrainLoop	mixed	San Francisco	~\$500K	Seed: ~\$500K, early/March 202...	Post-training research/product lab that turns enterprise data and evals into RL-fine-tuned ('reasoning fine-tuning') expert models for long-horizon d...

Bootstrapped / public (RL-env work, no qualifying venture round) — 2

- **AIChamp** (bootstrapped, San Francisco, CA...) — Builds custom RL environments and 'Virtual Gym' simulations for training and evaluating tool-using agents on ...
- **Innodata** (public, Ridgefield Park, ...) — Data-engineering incumbent offering custom RL gyms/environments as a managed service with deterministic evalu...

Also relevant but not enriched in this pass: **Surge AI** (bootstrapped; operates an internal RL-environments org but has taken no qualifying VC — its reported ~\$1B raise was still unclosed as of this session) and **Invisible Technologies** (bootstrapped/profitable; builds RL environments from enterprise workflows). Both belong in this "does the work, no qualifying VC" bucket.

Excluded — out of scope (11)

Surfaced by discovery but filtered out on inspection, with the reason:

- **Applied Compute** — inScope=false: core product is a vertically-integrated enterprise RL AGENT (customer-facing 'Specific Intelligence'), not a saleable RL environment /...
- **Cekura (formerly Vocera)** — Product is agent QA/testing/observability (simulated-user eval + production monitoring of deployed voice/chat agents), not RL environments/tasks/rewa...
- **Cua (trycua)** — Agent-sandbox/container infrastructure, not a core RL-environment/task/verifier product (though it ships benchmarks and is used for RL/eval), hence a...
- **Cyberdesk** — Scope=false: Cyberdesk is a virtual-desktop sandbox / computer-use-agent infrastructure and production RPA-style automation product, not an RL-enviro...
- **Daytona** — Sandbox/compute infrastructure provider, not a core RL-environment/task/verifier product — RL is a use case its sandboxes serve, hence adjacent-infra...
- **General Intuition** — OUT OF SCOPE: General Intuition is a frontier foundation-model lab building action/world models for embodied/spatial-temporal reasoning (gaming-to-ro...
- **Gensyn** — OUT OF SCOPE as an env/task/verifier vendor: Gensyn builds decentralized RL COMPUTE INFRASTRUCTURE (RL Swarm) for post-training, not RL environments,...
- **Hamming AI** — Same scope reasoning as Cekura: automated testing/simulation/monitoring/governance for deployed AI voice agents (eval + reliability), not RL environm...
- **Lucid** — Scope=false: Lucid's learned world-model 'RL gym' targets robotics (embodied training) and gaming, i.e. physical-world/embodied and game simulation, ...

- **Matrices** — Major scope/identity discrepancy: concrete public product evidence (Tracxn, Crunchbase, bizapedia) describes an AI-powered spreadsheet agent product,...
 - **TensorZero** — OUT OF SCOPE as an env/task/verifier product: TensorZero is an LLM OPTIMIZATION GATEWAY / LLMOps stack. It uses RL feedback loops (RLHF/DPO) to optim...
-

What the landscape shows (synthesis)

- **The category is real and crowded fast.** ~60 venture-funded companies now explicitly sell RL environments, tasks, or verifiers — most founded 2024–2026, most seed/pre-seed, heavily YC-originated (HUD, Refresh, Aviro, Idler, Maingen, Rubric AI, RunRL, Cyberdesk, Lucidic, TLDC, Halluminate, Osmosis).
- **Three sub-markets are forming:** (1) **environments/gyms** — the largest, computer-use & coding-agent focused (Fleet, Mechanize, DeepTune, Chakra, HUD, Refresh, Runloop, Plato); (2) **verifiers/RLVR & formal reward** — the fastest- scaling on capital (Harmonic ~\$875M val, Axiom Math \$64M seed, Judgment Labs \$32M, Gray Swan \$40M); (3) **task/expert-data foundries** adding RL-env lines (Mercor, Turing, Snorkel, Toloka, SuperAnnotate, iMerit, Micro1).
- **Biggest rounds ≠ pure-plays.** The heaviest capital sits with verifier/reasoning and world-model-adjacent plays; the environment pure-plays are mostly small seeds except Fleet (~725–750Mval) and DeepTune (43M Series A, a16z).
- **Consolidation already started:** Sepal → Mercor, OpenPipe → CoreWeave, Predibase → Rubrik — the pattern the Scale/Surge incumbents predicted.
- **Global but US-dominated:** UK (AgileRL, Huzzle, General Reasoning), France/CH (Andromede, Adaptive ML), NL (Toloka), India (Composio, Maxim, iMerit) present, but the center of gravity is SF.